

# Confusion Matrix

Manu Nellutla

8/26/2020

## Confusion Matrix

Measuring a model by accuracy can be misleading as the data used to train and test can be biased or have a uneven representation of the desired population mix.

To address this issue we often use **Confusion Matrix** to validate our models performance based on few additional metrics. But first a confusion matrix

$n = TP + TN + FP + FN$		Actual	
		Yes	No
Predicted	Yes	TP	FP
	No	FN	TN

## Intro

population (n) = all rows = (TP + FP + TN + FN)

- $Accuracy = \frac{(TP+TN)}{(TP+TN+FP+FN)}$  - proportion of correct classifications.
- $Precision = \frac{TP}{TP+FP}$  - This is also called 'Pos Prediction'
- $Recall = \frac{TP}{TP+FN}$  - This is also called 'Sensitivity'

## load data

```
library(dslabs)
data(heights)
summary(heights)
```

```
##      sex      height
## Female:238  Min.   :50.00
## Male   :812  1st Qu.:66.00
##                Median :68.50
##                Mean   :68.32
##                3rd Qu.:71.00
##                Max.   :82.68
```

## split dataset

Using **caret.createDataPartion** [<https://www.rdocumentation.org/packages/caret/versions/6.0-86/topics/createDataPartition>]  
(<https://www.rdocumentation.org/packages/caret/versions/6.0-86/topics/createDataPartition>)]

```
library("caret")
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
y = heights$sex
x = heights$height
set.seed(2, sample.kind = "Rounding")
```

```
## Warning in set.seed(2, sample.kind = "Rounding"): non-uniform 'Rounding' sampler
## used
```

```
# split at 0,5
test_index = createDataPartition(y, times = 1, p = 0.5, list=FALSE)

test_set = heights[test_index,]
train_set = heights[-test_index,]

summary(test_set)
```

```
##      sex      height
## Female:119  Min.    :50.00
## Male   :406  1st Qu.:66.00
##                Median :68.11
##                Mean    :68.36
##                3rd Qu.:71.00
##                Max.    :82.68
```

```
summary(train_set)
```

```
##      sex      height
## Female:119  Min.    :50.00
## Male   :406  1st Qu.:66.00
##                Median :68.90
##                Mean    :68.28
##                3rd Qu.:71.00
##                Max.    :80.00
```

## predic with best accuracy

Now that we know the best cutoff height is 64 lets predict with new cutoff

```
library(purrr)
```

```
##
## Attaching package: 'purrr'
```

```
## The following object is masked from 'package:caret':  
##  
## lift
```

```
y_hat <- ifelse(test_set$height > 64, "Male", "Female") %>% factor(levels = levels(test_set$sex))  
y_hat <- factor(y_hat)  
mean(y_hat == test_set$sex)
```

```
## [1] 0.8171429
```

## confusion matrix

```
confusionMatrix(data=y_hat, reference=test_set$sex)
```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction Female Male
##      Female      50   27
##      Male       69  379
##
##           Accuracy : 0.8171
##           95% CI : (0.7814, 0.8493)
##      No Information Rate : 0.7733
##      P-Value [Acc > NIR] : 0.008354
##
##           Kappa : 0.4041
##
##  McNemar's Test P-Value : 2.857e-05
##
##           Sensitivity : 0.42017
##           Specificity : 0.93350
##      Pos Pred Value : 0.64935
##      Neg Pred Value : 0.84598
##           Prevalence : 0.22667
##           Detection Rate : 0.09524
##      Detection Prevalence : 0.14667
##      Balanced Accuracy : 0.67683
##
##           'Positive' Class : Female
##

```

## Understanding results

If you look at the results we understand that **accuracy is about 81%**.

That looks good. But is it?

If you look at the “**Pos Prediction above it is 64%**”. This is our actual precision which tells us that proportion of ‘True Positives’ are low.

**Sensitivity/ Recall** is the proportion of correct classification for a class with in the population. > *In the above table*: predicted female = 50; actual female = 119 (look at test set summary)

Why is this? it is because of the cutoff we chose. I chose to classify male if height > 64. However, the mean of female height is 65. This makes the model intuitively wrong.

But 'Male' height cutoff was calculated as  $(\mu - 2 * sd)$  which is right based on our 95% CI.

What we can infer is that the **PREVALANCE** of the dataset, which means proportion of one class within the population, is uneven..

## Maximize F1 score

One preferred metric is **balanced accuracy**. Because specificity and sensitivity are rates, it is more appropriate to compute the harmonic average. In fact, the F1-score, a widely used one-number summary, is the harmonic average of precision and recall.

$$F1 = \frac{2 * [precision * recall]}{[precision + recall]}$$

```
# maximize F-score
cutoff <- seq(61, 70)
F_1 <- map_dbl(cutoff, function(x){
  y_hat <- ifelse(train_set$height > x, "Male", "Female") %>%
    factor(levels = levels(test_set$sex))
  F_meas(data = y_hat, reference = factor(train_set$sex))
})

F_1
```

```
## [1] 0.2112676 0.3225806 0.4393064 0.5656566 0.6052632 0.6142322 0.5820433
## [8] 0.5573770 0.5080831 0.4679089
```


```
max(F_1)
```

```
## [1] 0.6142322
```

For the same cutoff we find F1 score using “**F\_means()**” function of *caret* package. Max F1 score is around 61%.

## lets plot of F1 scores

```
data.frame(cutoff, F_1) %>%  
  ggplot(aes(cutoff, F_1)) +  
    geom_point() +  
    geom_line()
```



```
max(F_1)
```

```
## [1] 0.6142322
```

```
best_cutoff <- cutoff[which.max(F_1)]  
best_cutoff
```

```
## [1] 66
```

From the R console we see the best cutoff is not at 66 and not 64. This is better as the mean(female\_heights) is 65.

## lets check our metrics again for best cutoff

```
y_hat <- ifelse(test_set$height > best_cutoff, "Male", "Female") %>%  
  factor(levels = levels(test_set$sex))  
sensitivity(data = y_hat, reference = test_set$sex)
```

```
## [1] 0.6806723
```

```
specificity(data = y_hat, reference = test_set$sex)
```

```
## [1] 0.8349754
```

what about confusion matrix

```
confusionMatrix(data=y_hat, reference=test_set$sex)
```



```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction Female Male
##      Female      81   67
##      Male       38  339
##
##           Accuracy : 0.8
##           95% CI : (0.7632, 0.8334)
##      No Information Rate : 0.7733
##      P-Value [Acc > NIR] : 0.078192
##
##           Kappa : 0.4748
##
##  McNemar's Test P-Value : 0.006285
##
##           Sensitivity : 0.6807
##           Specificity : 0.8350
##      Pos Pred Value : 0.5473
##      Neg Pred Value : 0.8992
##           Prevalence : 0.2267
##           Detection Rate : 0.1543
##      Detection Prevalence : 0.2819
##      Balanced Accuracy : 0.7578
##
##           'Positive' Class : Female
##

```

Lets compare before and after

Without F1	With F1
------------	---------

### Confusion Matrix and Statistics

	Reference	
Prediction	Female	Male
Female	50	27
Male	69	379

Accuracy : 0.8171  
95% CI : (0.7814, 0.8493)  
No Information Rate : 0.7733  
P-Value [Acc > NIR] : 0.008354

Kappa : 0.4041

McNemar's Test P-Value : 2.857e-05

Sensitivity : 0.42017  
Specificity : 0.93350  
Pos Pred Value : 0.64935  
Neg Pred Value : 0.84598  
Prevalence : 0.22667  
Detection Rate : 0.09524  
Detection Prevalence : 0.14667  
Balanced Accuracy : 0.67683

'Positive' Class : Female

### Confusion Matrix and Statistics

	Reference	
Prediction	Female	Male
Female	81	67
Male	38	339

Accuracy : 0.8  
95% CI : (0.7632, 0.8334)  
No Information Rate : 0.7733  
P-Value [Acc > NIR] : 0.078192

Kappa : 0.4748

McNemar's Test P-Value : 0.006285

Sensitivity : 0.6807  
Specificity : 0.8350  
Pos Pred Value : 0.5473  
Neg Pred Value : 0.8992  
Prevalence : 0.2267  
Detection Rate : 0.1543  
Detection Prevalence : 0.2819  
Balanced Accuracy : 0.7578

'Positive' Class : Female

## ROC Curves

```
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':  
##  
## filter, lag
```


```
## The following objects are masked from 'package:base':  
##  
## intersect, setdiff, setequal, union
```

```
cutoffs <- c(50, seq(60, 75), 80)  
height_cutoff <- map_df(cutoffs, function(x){  
  y_hat <- ifelse(test_set$height > x, "Male", "Female") %>%  
    factor(levels = c("Female", "Male"))  
  list(method = "Height cutoff",  
        FPR = 1-specificity(y_hat, test_set$sex),  
        TPR = sensitivity(y_hat, test_set$sex))  
})
```

## Plotting the ROC curve


A widely used plot that does this is the **receiver operating characteristic (ROC)** curve. The ROC curve plots *sensitivity (TPR)* versus *1 - specificity or the false positive rate (FPR)*.

```
# plot both curves together  
height_cutoff %>%  
  ggplot(aes(FPR, TPR, color = method)) +  
  geom_line() +  
  geom_point() +  
  xlab("1 - Specificity") +  
  ylab("Sensitivity")
```



another plot

```
library(ggrepel)
map_df(cutoffs, function(x){
  y_hat <- ifelse(test_set$height > x, "Male", "Female") %>%
    factor(levels = c("Female", "Male"))
  list(method = "Height cutoff",
        cutoff = x,
        FPR = 1-specificity(y_hat, test_set$sex),
        TPR = sensitivity(y_hat, test_set$sex))
}) %>%
  ggplot(aes(FPR, TPR, label = cutoff)) +
  geom_line() +
  geom_point() +
  geom_text_repel(nudge_x = 0.01, nudge_y = -0.01)
```



## Precision-Recall plot.

However, ROC curves have one weakness and it is that neither of the measures plotted depend on prevalence. In cases in which prevalence matters, we may instead make a precision-recall plot, which has a similar idea with ROC curve.

```

probs <- seq(0, 1, length.out = 10)
guessing <- map_df(probs, function(p){
  y_hat <- sample(c("Male", "Female"), length(test_index),
    replace = TRUE, prob=c(p, 1-p)) %>%
    factor(levels = c("Female", "Male"))
  list(method = "Guess",
    recall = sensitivity(y_hat, test_set$sex),
    precision = precision(y_hat, test_set$sex))
})


height_cutoff <- map_df(cutoffs, function(x){
  y_hat <- ifelse(test_set$height > x, "Male", "Female") %>%
    factor(levels = c("Female", "Male"))
  list(method = "Height cutoff",
    recall = sensitivity(y_hat, test_set$sex),
    precision = precision(y_hat, test_set$sex))
})

bind_rows(guessing, height_cutoff) %>%
  ggplot(aes(recall, precision, color = method)) +
  geom_line() +
  geom_point()

```

```
## Warning: Removed 1 row(s) containing missing values (geom_path).
```

```
## Warning: Removed 1 rows containing missing values (geom_point).
```



## with releval

Meaning ... If we change positives to mean male instead of females.



```
guessing <- map_df(probs, function(p){
  y_hat <- sample(c("Male", "Female"), length(test_index), replace = TRUE,
    prob=c(p, 1-p)) %>%
    factor(levels = c("Male", "Female"))
  list(method = "Guess",
    recall = sensitivity(y_hat, relevel(test_set$sex, "Male", "Female")),
    precision = precision(y_hat, relevel(test_set$sex, "Male", "Female")))
})

height_cutoff <- map_df(cutoffs, function(x){
  y_hat <- ifelse(test_set$height > x, "Male", "Female") %>%
    factor(levels = c("Male", "Female"))
  list(method = "Height cutoff",
    recall = sensitivity(y_hat, relevel(test_set$sex, "Male", "Female")),
    precision = precision(y_hat, relevel(test_set$sex, "Male", "Female")))
})

bind_rows(guessing, height_cutoff) %>%
  ggplot(aes(recall, precision, color = method)) +
  geom_line() +
  geom_point()
```

```
## Warning: Removed 1 row(s) containing missing values (geom_path).
```

```
## Warning: Removed 1 rows containing missing values (geom_point).
```

